

ADENINE — A Data Exploration pIpeLINE

Samuele Fiorini
Federico Tomasi
Annalisa Barla

SAMUELE.FIORINI@DIBRIS.UNIGE.IT
FEDERICO.TOMASI@DIBRIS.UNIGE.IT
ANNALISA.BARLA@UNIGE.IT

*Department of Informatics, Bioengineering,
Robotics and System Engineering (DIBRIS)
University of Genoa
Genoa, I-16146, Italy*

Editor: Editor name

Abstract

In this paper we introduce **Adenine**, a **Python** framework for data exploration. The main goal of **Adenine**, is twofold: helping researchers and data scientists achieving a first and quick overview on the main structures underlying their data and choosing the most suitable unsupervised learning pipeline for the problem at hand. This software tool encompasses state of the art techniques for: missing values imputing, preprocessing, dimensionality reduction and clustering tasks. **Adenine** exploits both process and thread level parallelism and it is capable of generating nice and clean publication-ready figures along with a quantitative description of the algorithms performance. **Adenine** is released under FreeBSD license and it can be downloaded from <http://slipguru.github.io/adenine/>.

Keywords: Data exploration, unsupervised learning, RNA-Seq gene expression

1. Introduction

Data exploration can be an extremely helpful starting point for many data analysis tasks. Researchers and data scientists are often asked to extract meaningful information from collections of complex and possibly high-dimensional data coming from heterogenous contexts. For instance, in biomedical scenarios, physicians are likely to be interested in answering some biological questions starting from a set of observations collected from a pool of subjects enrolled in a study. Possible investigations can be: *is there any relevant stratification among subjects?* or *is it possible to discriminate between cases and controls from my observations?*. Starting from a given dataset, the information needed to answer such questions may be immediate, non-trivial to extract or even completely absent. In these situations, a preliminary data exploration step can be not only good practice, but also a mandatory starting point for further investigations. In this context, several machine learning and data mining techniques were developed over the years. Among those we focus on the four most popular families: (i) missing values imputing, (ii) data preprocessing, (iii) dimensionality reduction and (iv) unsupervised clustering.

In the last few years, a fair number of data exploration software tools and libraries were released. At a very coarse grain we can group them into two sets: interactive GUI-based and non-interactive script-based applications. Of the first group, we recall *Divvy* (Lewis

et al., 2013), a software tool that performs dimensionality reduction and clustering on input datasets. *Divvy* is a light framework, although it is only Mac OS X specific, it heavily lacks in documentation, its collection of C/C++ algorithm implementations does not cover common strategies such as Kernel Principal Component Analysis (KPCA) (Schölkopf et al., 1997) or Hierarchical Clustering (Friedman et al., 2001) and it does not offer strategies to perform automatic discoveries of the number of clusters. The most notable project that spans between the two families is *Orange* (Demšar et al., 2013), a data mining software suite that offers both visual programming front-end and standard `Python` APIs. In the context of data exploration, *Orange* can be successfully employed. However it does not support automatic pipeline generation, hence it requires the user to manually create each pipeline. On the other hand, as of today *Orange* lacks in several nonlinear methods such as Spectral Embedding (Ng et al., 2002), Locally Linear Embedding (Roweis and Saul, 2000) or Spectral Clustering (Shi and Malik, 2000).

We developed **Adenine**, a script-based `Python` tool for data exploration that, starting from a set of predefined unsupervised algorithms, saves textual and graphical reports of an arbitrary large number of pipelines. Data imputing, preprocessing, dimensionality reduction and clustering strategies are here seen as building blocks. The user is only required to specify which blocks should be activated and **Adenine** takes care of automatically generate and run the pipelines made by all possible combinations of the selected algorithms. All **Adenine** algorithm implementation is inherited, or extended, from `scikit-learn` (Pedregosa et al., 2011) which is, to the best of our knowledge, the most complete machine learning open source library freely available online.

2. Implementation

From an implementative standpoint, **Adenine** is built around the concept of *pipeline*, that is a sequence of the four fundamental steps mentioned in Section 1 (see Figure 1).

For each step, a fair number of off-the-shelf algorithms are available (see Table 2).

In order to perform exploratory analysis of large datasets, **Adenine** can take advantage of different parallel computing paradigms. For instance, its pipelines are designed to be independent from each other, therefore they all run in parallel as separate `Python` processes on different cores. Moreover, since **Adenine** makes large use of `numpy` and `scipy`, it automatically benefits from their bindings with optimized linear algebra libraries (such as OpenBLAS¹, or Intel[®] MKL).

3. Experiments and results

To assess the quality of the obtained results, we tested **Adenine** on a set of synthetic and real dataset.

{parla qui dei test synth} {TGCA}

Table 1: Pipelines building blocks and relative references (which are not reported when the definition is given in Section 2).

Step	Algorithms	Ref.
Imputing	Mean	(Troyanskaya et al., 2001)
	Median	
	KNN	
Preprocessing	Recentring	
	Standardize	
	Normalize	
	MinMax	
Dimensionality reduction	Principal component Analysis (PCA)	(Jolliffe, 2002)
	Incremental PCA	(Ross et al., 2008)
	Randomized PCA	(Halko et al., 2011)
	Kernel PCA	(Schölkopf et al., 1997)
	Isomap	(Tenenbaum et al., 2000)
	Locally Linear Embedding	(Roweis and Saul, 2000)
	Spectral Embedding	(Ng et al., 2002)
	Multidimensional Scaling	(Borg and Groenen, 2005)
	t-Distributed Stochastic Neighbor Embedding (t-SNE)	(Van der Maaten and Hinton, 2008)
Clustering	K-means	(Bishop, 2006)
	Affinity propagation	(Frey and Dueck, 2007)
	Mean Shift	(Comaniciu and Meer, 2002)
	Spectral	(Shi and Malik, 2000)
	Hierarchical	(Friedman et al., 2001)

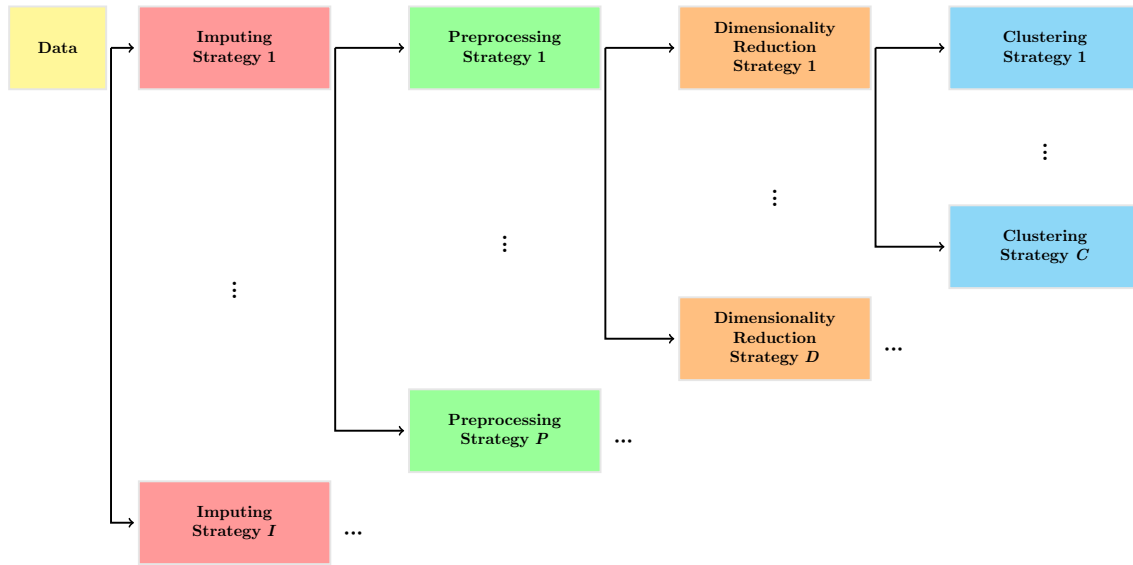


Figure 1: A schematic representation of the **Adenine** workflow. The list of building blocks available for each step is summarized in Table 1.

4. Conclusions

References

- Christopher M Bishop. Pattern recognition. *Machine Learning*, 2006.
- Ingwer Borg and Patrick JF Groenen. *Modern multidimensional scaling: Theory and applications*. Springer Science & Business Media, 2005.
- Dorin Comaniciu and Peter Meer. Mean shift: A robust approach toward feature space analysis. *Pattern Analysis and Machine Intelligence, IEEE Transactions on*, 24(5):603–619, 2002.
- Janez Demšar, Tomaž Curk, Aleš Erjavec, Črt Gorup, Tomaž Hočevar, Mitar Milutinović, Martin Možina, Matija Polajnar, Marko Toplak, Anže Starič, et al. Orange: data mining toolbox in python. *The Journal of Machine Learning Research*, 14(1):2349–2353, 2013.
- Brendan J Frey and Delbert Dueck. Clustering by passing messages between data points. *science*, 315(5814):972–976, 2007.
- Jerome Friedman, Trevor Hastie, and Robert Tibshirani. *The elements of statistical learning*, volume 1. Springer series in statistics Springer, Berlin, 2001.

1. <http://www.openblas.net/>

- Nathan Halko, Per-Gunnar Martinsson, and Joel A Tropp. Finding structure with randomness: Probabilistic algorithms for constructing approximate matrix decompositions. *SIAM review*, 53(2):217–288, 2011.
- Ian Jolliffe. *Principal component analysis*. Wiley Online Library, 2002.
- Joshua M Lewis, Virginia R De Sa, and Laurens Van Der Maaten. Divvy: fast and intuitive exploratory data analysis. *The Journal of Machine Learning Research*, 14(1):3159–3163, 2013.
- Andrew Y Ng, Michael I Jordan, Yair Weiss, et al. On spectral clustering: Analysis and an algorithm. *Advances in neural information processing systems*, 2:849–856, 2002.
- F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830, 2011.
- David A Ross, Jongwoo Lim, Ruei-Sung Lin, and Ming-Hsuan Yang. Incremental learning for robust visual tracking. *International Journal of Computer Vision*, 77(1-3):125–141, 2008.
- Sam T Roweis and Lawrence K Saul. Nonlinear dimensionality reduction by locally linear embedding. *Science*, 290(5500):2323–2326, 2000.
- Bernhard Schölkopf, Alexander Smola, and Klaus-Robert Müller. Kernel principal component analysis. In *Artificial Neural Networks ICANN’97*, pages 583–588. Springer, 1997.
- Jianbo Shi and Jitendra Malik. Normalized cuts and image segmentation. *Pattern Analysis and Machine Intelligence, IEEE Transactions on*, 22(8):888–905, 2000.
- Joshua B Tenenbaum, Vin De Silva, and John C Langford. A global geometric framework for nonlinear dimensionality reduction. *science*, 290(5500):2319–2323, 2000.
- Olga Troyanskaya, Michael Cantor, Gavin Sherlock, Pat Brown, Trevor Hastie, Robert Tibshirani, David Botstein, and Russ B Altman. Missing value estimation methods for dna microarrays. *Bioinformatics*, 17(6):520–525, 2001.
- Laurens Van der Maaten and Geoffrey Hinton. Visualizing data using t-sne. *Journal of Machine Learning Research*, 9(2579-2605):85, 2008.